

# Database Connection Pooling in Django

## 1. Introduction

Database connection pooling is a technique to reuse existing database connections instead of creating new ones for every request. This helps reduce connection overhead and improves performance.

Django supports connection pooling via the `CONN_MAX_AGE` setting, which controls how long a database connection remains open before being closed.

---

## 2. How Django Handles Connection Pooling

By default, Django creates a **new database connection** for each request and closes it after the request is completed. This can lead to performance issues when handling multiple concurrent requests.

To enable connection pooling, we use:

```
DATABASES = {  
    'default': {  
        'ENGINE': 'django.db.backends.postgresql',  
        'NAME': 'mydb',  
        'USER': 'myuser',  
        'PASSWORD': 'mypassword',  
        'HOST': 'localhost',  
        'PORT': '5432',  
        'CONN_MAX_AGE': 600, # Keep connections open for 10 minutes  
    }  
}
```

# Database Pooling

## Behavior of `CONN_MAX_AGE`

1. If a connection is **available and within the timeout**, Django **reuses it**.
  2. If a connection is **idle for more than 10 minutes**, Django **closes it**.
  3. If a new request comes after 10 minutes, Django **creates a new connection**.
- 

## 3. Does Pooling Require Query Modifications?

No! Django automatically manages connection reuse when `CONN_MAX_AGE` is set. Example:

```
user = User.objects.get(id=1) # This query will reuse an open connection if available
```

You don't need to manually handle pooling in your queries.

---

## 4. Database Connection Limits

Every database has a limit on the maximum number of concurrent connections:

- PostgreSQL (`max_connections`): Default **100** (can be increased via `ALTER SYSTEM SET max_connections = 500;`)
- MySQL (`max_connections`): Default **151** (can be changed in `my.cnf`)

If too many connections are created, the database may reject new connections, leading to errors.

---

## 5. Optimizing Pooling with PgBouncer

## Database Pooling

If your application has **high concurrent traffic**, using PgBouncer can further optimize pooling. PgBouncer allows many users to share a small number of actual database connections.

### PgBouncer Configuration (**pgbouncer.ini**)

```
[pgbouncer]
max_client_conn = 1000
default_pool_size = 50
pool_mode = transaction # Best for Django
```

### Update Django Settings for PgBouncer

```
DATABASES = {
    'default': {
        'ENGINE': 'django.db.backends.postgresql',
        'NAME': 'mydb',
        'USER': 'myuser',
        'PASSWORD': 'mypassword',
        'HOST': '127.0.0.1', # PgBouncer host
        'PORT': '6432', # PgBouncer port
        'CONN_MAX_AGE': 600,
    }
}
```

### PgBouncer Benefits

- ☒ Reduces database connection overhead
- ☒ Supports **1000+ users** with only **50 real DB connections**
- ☒ Prevents hitting **max\_connections** limit

---

## 6. Best Practices

## Database Pooling

✓ Use a single database user for Django {

```
'default': {  
  
    'ENGINE': 'django.db.backends.postgresql',  
  
    'NAME': 'mydb',  
  
    'USER': 'myuser',  
  
    'PASSWORD': 'mypassword',  
  
    'HOST': '127.0.0.1', # PgBouncer host  
  
    'PORT': '6432', # PgBouncer port  
  
    'CONN_MAX_AGE': 600,  
  
}
```

(myuser) instead of separate users per request. ✓ Increase **max\_connections** if your database supports high traffic. ✓ Enable **PgBouncer** for production environments. ✓ Monitor active connections with:

```
SELECT * FROM pg_stat_activity;
```

---

## 7. Conclusion

- Django's **CONN\_MAX\_AGE** allows automatic connection pooling.
- **PgBouncer** improves scalability by reducing the number of actual DB connections.
- No need to modify queries, as Django manages connections automatically.
- For high concurrency, use **PgBouncer** + **optimized DB settings**.